

+ Code

+ Text

21BIT0389 Pavithra S Assingment 4

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv('/content/winequality-red.csv')
df
```




	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	6
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	5
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	6

1599 rows × 12 columns



```
plt.pie(df.quality)  
plt.show
```

```
<function matplotlib.pyplot.show(close=None, block=None)>
```



```
sns.distplot(df.quality)
```

```
<ipython-input-4-e8684199aa87>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see


<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.quality)  
<Axes: xlabel='quality', ylabel='Density'>
```




```
plt.bar(df.quality,df.alcohol)
```

<BarContainer object of 1599 artists>



```
sns.jointplot(x='quality',y='pH',data=df)
```

<seaborn.axisgrid.JointGrid at 0x7c9175c7bd60>



sns.heatmap(df.corr())

<Axes: >

```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7c916a514a30>
```

```
df.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	dens
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000	1.003



```
df.isnull().any()
```

```
fixed acidity      False
volatile acidity   False
citric acid        False
residual sugar     False
chlorides          False
free sulfur dioxide False
total sulfur dioxide False
density            False
pH                 False
sulphates          False
alcohol            False
quality            False
dtype: bool
```

```
df.median()
```

```
fixed acidity      7.90000  
volatile acidity   0.52000  
citric acid        0.26000  
residual sugar     2.20000  
chlorides          0.07900  
free sulfur dioxide 14.00000  
total sulfur dioxide 38.00000  
density           0.99675  
pH                3.31000  
sulphates          0.62000  
alcohol           10.20000  
quality           6.00000  
dtype: float64
```

```
sns.boxplot(df['fixed acidity'])
```



<Axes: >

```

-- |
q1=df['fixed acidity'].quantile(0.25)
q3=df['fixed acidity'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+(1.5*IQR)
lower_limit=q1-(1.5*IQR)
12 |
df['fixed acidity']=np.where(df['fixed acidity']>upper_limit,7.9,df['fixed acidity']
|
|
|
sns.boxplot(df['fixed acidity'])


```

<Axes: >



```
sns.boxplot(df['volatile acidity'])
```

<Axes: >




```
q1=df['volatile acidity'].quantile(0.25)
q3=df['volatile acidity'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+(1.5*IQR)
lower_limit=q1-(1.5*IQR)
```

```
df['volatile acidity']=np.where(df['volatile acidity']>upper_limit,0.52,df['volatil
```

```
sns.boxplot(df['volatile acidity'])
```

<Axes: >



```
sns.boxplot(df['citric acid'])
```

Statistic	Value
Minimum	0.78
Q1	0.78
Median	0.78
Q3	0.80
Maximum	0.80
Outliers	1.0


```
acid'] = np.where(df['citric acid']>upper_limit,0.26,df['citric acid'])
```

<Axes: >



```
sns.boxplot(df['residual sugar'])
```

<Axes: >




```
q1=df['residual sugar'].quantile(0.25)
q3=df['residual sugar'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df['residual sugar'] = np.where(df['residual sugar']>upper_limit,2.2,df['residual sugar'])
```

```
sns.boxplot(df['residual sugar'])
```

<Axes: >




```
q1=df['residual sugar'].quantile(0.25)
q3=df['residual sugar'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df['residual sugar'] = np.where(df['residual sugar']>upper_limit,2.2,df['residual sugar'])
```

```
sns.boxplot(df['residual sugar'])
```

<Axes: >




```
q1=df['residual sugar'].quantile(0.25)
q3=df['residual sugar'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df['residual sugar'] = np.where(df['residual sugar']<lower_limit,2.2,df['residual sugar'])
```

```
sns.boxplot(df['residual sugar'])
```


<Axes: >



```
q1=df['residual sugar'].quantile(0.25)
q3=df['residual sugar'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df['residual sugar'] = np.where(df['residual sugar']>upper_limit,2.2,df['residual sugar'])
```

```
sns.boxplot(df['residual sugar'])
```

<Axes: >



```
q1=df['chlorides'].quantile(0.25)
```

```
q3=df['chlorides'].quantile(0.75)
```

```
IQR=q3-q1
```

```
upper_limit=q3+1.5*IQR
```

```
lower_limit=q1-1.5*IQR
```

|

■

|

```
df['chlorides'] = np.where(df['chlorides']>upper_limit,0.08,df['chlorides'])
```

|

■

|

```
sns.boxplot(df['chlorides'])
```

<Axes: >

0.12

```

q1=df['chlorides'].quantile(0.25)
q3=df['chlorides'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR

```

```
df['chlorides'] = np.where(df['chlorides']>upper_limit,0.08,df['chlorides'])
```

0.06

```
sns.boxplot(df['chlorides'])
```

<Axes: >

0.10

0.08

0.06

0.04

0.02


0

```
q1=df['chlorides'].quantile(0.25)
q3=df['chlorides'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df['chlorides'] = np.where(df['chlorides']<lower_limit,0.08,df['chlorides'])
```

```
sns.boxplot(df['chlorides'])
```

<Axes: >



```
q1=df['chlorides'].quantile(0.25)
q3=df['chlorides'].quantile(0.75)
IQR=q3-q1
```


```
upper_limit=q3+1.5*IQR
```

```
lower_limit=q1-1.5*IQR
```

```
df['chlorides'] = np.where(df['chlorides']>upper_limit,0.08,df['chlorides'])
```


```
sns.boxplot(df['chlorides'])
```

<Axes: >



```
sns.boxplot(df['free sulfur dioxide'])
```

<Axes: >



```
q1=df['free sulfur dioxide'].quantile(0.25)
```

```
q3=df['free sulfur dioxide'].quantile(0.75)
```

```
IQR=q3-q1
```


```
upper_limit=q3+1.5*IQR
```

```
lower_limit=q1-1.5*IQR
```

```
df['free sulfur dioxide'] = np.where(df['free sulfur dioxide']>upper_limit,14,df['free sulfur dioxide'])
```

```
sns.boxplot(df['free sulfur dioxide'])
```

<Axes: >



```
sns.boxplot(df['total sulfur dioxide'])
```

<Axes: >

300

```
q1=df['total sulfur dioxide'].quantile(0.25)
```

```
q3=df['total sulfur dioxide'].quantile(0.75)
```

```
IQR=q3-q1
```


```
upper_limit=q3+1.5*IQR
```

```
lower_limit=q1-1.5*IQR
```

```
df['total sulfur dioxide'] = np.where(df['total sulfur dioxide']>upper_limit,38,df['total sulfur dioxide'])
```

```
sns.boxplot(df['total sulfur dioxide'])
```

<Axes: >





```
q1=df['total sulfur dioxide'].quantile(0.25)
q3=df['total sulfur dioxide'].quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df['total sulfur dioxide'] = np.where(df['total sulfur dioxide']>upper_limit,38,df['total sulfur dioxide'])
```


```
sns.boxplot(df['total sulfur dioxide'])
```

<Axes: >



```
sns.boxplot(df.density)
```

<Axes: >




```
q1=df.density.quantile(0.25)
q3=df.density.quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df.density = np.where(df.density>upper_limit,1,df.density)
```

```
sns.boxplot(df.density)
```


<Axes: >



```
q1=df.density.quantile(0.25)
q3=df.density.quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df.density = np.where(df.density<lower_limit,1,df.density)
```

```
sns.boxplot(df.density)
```



<Axes: >

4.0

```
q1=df.pH.quantile(0.25)
```

```
q3=df.pH.quantile(0.75)
```

```
IQR=q3-q1
```

```
upper_limit=q3+1.5*IQR
```

```
lower_limit=q1-1.5*IQR
```

|

|

|

```
df.pH = np.where(df.pH<lower_limit,3.3,df.pH)
```

|



|

```
sns.boxplot(df.pH)
```


```
<Axes: >
```

```
q1=df.pH.quantile(0.25)
q3=df.pH.quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df.pH = np.where(df.pH>upper_limit,3.3,df.pH)
```


```
sns.boxplot(df.pH)
```

```
<Axes: >
```



```
sns.boxplot(df.sulphates)
```

<Axes: >




```
q1=df.sulphates.quantile(0.25)
q3=df.sulphates.quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df.sulphates = np.where(df.sulphates>upper_limit,0.62,df.sulphates)
```

```
sns.boxplot(df.sulphates)
```

<Axes: >




```
q1=df.sulphates.quantile(0.25)
q3=df.sulphates.quantile(0.75)
IQR=q3-q1
upper_limit=q3+1.5*IQR
lower_limit=q1-1.5*IQR
```

```
df.sulphates = np.where(df.sulphates>upper_limit,0.62,df.sulphates)
```

```
sns.boxplot(df.sulphates)
```


<Axes: >



```
sns.boxplot(df.alcohol)
```

<Axes: >

15



```
q1=df.alcohol.quantile(0.25)
```

```
q3=df.alcohol.quantile(0.75)
```

```
IQR=q3-q1
```

```
upper_limit=q3+1.5*IQR
```

```
lower_limit=q1-1.5*IQR
```

14

|

|

```
df.alcohol = np.where(df.alcohol>upper_limit,10.2,df.alcohol)
```

11




|

```
sns.boxplot(df.alcohol)
```

<Axes: >

```
sns.boxplot(df.quality)
```

<Axes: >



```
df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1599 entries, 0 to 1598

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64

```

4 chlorides          1599 non-null    float64
5 free sulfur dioxide 1599 non-null    float64
6 total sulfur dioxide 1599 non-null    float64
7 density            1599 non-null    float64
8 pH                 1599 non-null    float64
9 sulphates          1599 non-null    float64
10 alcohol           1599 non-null    float64
11 quality           1599 non-null    int64

```

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

```
x=df.drop(columns=['quality'],axis=1)
```

```
x.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

```
y=df.quality
```

```
y.head()
```

```

0    5
1    5
2    5
3    6
4    5

```

```
Name: quality, dtype: int64
```

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

scale=MinMaxScaler()

x_s=pd.DataFrame(scale.fit_transform(x),columns=x.columns)

x_train,x_test,y_train,y_test=train_test_split(x_s,y,test_size=0.2,random_state=0)

x_train.shape,y_train.shape

((1279, 11), (1279,))

x_test.shape,y_test.shape

((320, 11), (320,))
```

Random Forest

```
from sklearn.ensemble import RandomForestClassifier

model=RandomForestClassifier(n_estimators=400)

model.fit(x_train,y_train)
```

▼ RandomForestClassifier

RandomForestClassifier(n_estimators=400)

```
y_pred=model.predict(x_test)
```

```
y_pred_train=model.predict(x_train)
```

```
from sklearn.metrics import accuracy_score
print(accuracy_score(y_train,y_pred_train))
```

```
1.0
```

```
print(accuracy_score(y_test,y_pred))
```

```
0.728125
```

```
from sklearn.metrics import classification_report,confusion_matrix
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	2
4	0.00	0.00	0.00	11
5	0.78	0.80	0.79	135
6	0.73	0.77	0.75	142
7	0.53	0.59	0.56	27
8	0.00	0.00	0.00	3
accuracy			0.73	320
macro avg	0.34	0.36	0.35	320
weighted avg	0.70	0.73	0.71	320


```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
```

```
confusion_matrix(y_test, y_pred)
```

```
array([[ 0,  0,  1,  1,  0,  0],  
       [ 0,  0,  6,  5,  0,  0],  
       [ 0,  0, 108, 25,  2,  0],  
       [ 0,  0, 23, 109, 10,  0],  
       [ 0,  0,  1,  8, 16,  2],  
       [ 0,  0,  0,  1,  2,  0]])
```


```
model.predict([[7.5, 0.8, 0.0, 2, 0.75, 15, 38, 0.96, 3.5, 0.5, 9.8]])
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names  
warnings.warn(  
array([5])
```



```
model.predict([[7.4, 0.7, 0.1, 2, 0.75, 14, 35, 0.96, 3.5, 0.5, 9.9]])
```

```
base.py:439: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names
```



```
model.predict([[7.9, 0.52, 0.26, 2.2, 0.08, 14, 38, 0.99, 3.3, 0.62, 10.2]])
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names  
warnings.warn(  
array([7])
```

